C-Suite AI Playbook

7 Proven MVPs That Delivered 6-Figure ROI in <30 Days

A Framework for Executive Teams Building AI Velocity

Introduction

This playbook distills 5+ years of experience directing AI programs for executives who have strong engineering teams but zero AI velocity. Inside, you'll find the 7 AI projects that consistently deliver measurable ROI in 10–16 weeks, the hiring roadmap that actually works, and the decision framework we use to choose between build vs. buy.

1. The 7 MVPs That Move the Needle

Conversational AI Agents for Customer Service
Phone agents that replace legacy IVRs. Handle 70–80% of inbound calls without human touch. ROI: $40K–$200K/month in labor savings + improved resolution rates.

Data Integrity & Reporting Pipelines
Zero-loss data flows from source systems to reporting dashboards. Board-ready ROI visibility. ROI: $25K–$100K in prevented data losses + reduced reporting overhead.

Computer Vision for Operations
Automated visual inspection, geospatial analysis, or asset tracking. Deploy at scale in 12–16 weeks. ROI: $50K–$300K in efficiency gains or compliance automation.

Intelligent Document Processing
Extract, classify, and route documents automatically. Cut manual processing by 60–80%. ROI: $30K–$150K in labor savings.

Predictive Analytics for Revenue/Risk
Forecasting models for churn, lifetime value, or operational risk. Inform strategy at board level. ROI: $50K–$500K depending on business impact.

Personalization Engines
Recommend products, content, or actions based on user behavior. Increase conversion 15–30%. ROI: $100K–$1M depending on transaction volume.

Search & Knowledge Retrieval (RAG)
Internal knowledge base with AI-powered search. Reduce support tickets 20–40%. ROI: $20K–$100K in labor savings.

2. Build vs. Buy: The Decision Framework

Build If:
• You have competitive differentiation (rare)
• Your problem is truly unique to your business
• You have 3–5 strong ML/data engineers
• Timeline is flexible (16+ weeks)

Buy If:
• Your problem is common (customer service, document processing, analytics)
• You need results in 4–12 weeks
• You lack in-house ML depth
• You want to avoid ongoing maintenance

3. The Hiring Roadmap for 10–16 Week Delivery

Immediate (Weeks 1–2):
Hire 1 senior AI architect (or director) to own the program. Non-negotiable.

Phase 1 (Weeks 2–4):
Hire 2–3 ML engineers or data engineers depending on your chosen MVP.

Phase 2 (Weeks 4–8):
Hire 1 DevOps/MLOps engineer to handle deployment, monitoring, and infrastructure.

Phase 3 (Weeks 8–16):
Hire a product-focused role (product manager or solutions engineer) to iterate based on user feedback.

4. What Success Looks Like

Week 4: Working prototype. Users can interact with it. Metrics baseline is set.
Week 8: Beta deployment. Real data flowing. ROI projections locked in.
Week 12: Production launch. Monitoring in place. Teams trained.
Week 16: Optimization phase. Gathering feedback. Planning Phase 2.

5. 5 Pitfalls That Kill AI Programs

1. **Hiring ML engineers before defining the problem.** You'll waste 8–12 weeks on wrong approaches.
2. **Picking a vendor before proving the problem is real.** Avoid lock-in; prototype first.
3. **Underestimating data prep work.** 60% of time goes to cleaning data, not modeling.
4. **Skipping MLOps infrastructure.** A beautiful model that can't scale is worthless.
5. **No executive sponsorship.** This is critical. Budget, headcount, and decisions all depend on it.

Ready to move from zero AI velocity to 10–16 week delivery?
Book a 30-minute diagnostic to discuss your specific use case and get a custom roadmap.